

CouchDB

an introduction

Play Along

<http://www.couchone.com/get>

<https://cloudant.com/>

<http://www.couchone.com/get>

<http://guide.couchdb.org/>

<http://wiki.apache.org/couchdb/>

<http://www.couchone.com/get>

- NoSQL
- JSON Documents
- HTTP RESTful API
- Incremental Map/Reduce
- Durable, Crash-Only Design

<http://www.couchone.com/get>

- Erlang
- Replication / Conflict Resolution
- Data stored in B+Tree
- Single file for each database
- Database is never locked
- Writes are serialized

Futon


Overview

+ Create Database ...

| Name | Size | Number of Documents | Update Seq |
|--------|---------|---------------------|------------|
| jibjab | 1.0 MB | 61 | 79 |
| _users | 20.1 KB | 2 | 5 |

Showing 1-2 of 2 databases

← Previous Page | Rows per page: 10 | Next Page →



CouchDB
relax

Tools

- Overview
- Configuration
- Replicator
- Status
- Test Suite

Recent Databases

- _users
- jibjab

Signup or Login

Futon on Apache CouchDB 1.0.1

http://127.0.0.1:5984/_utils

Basic API


```
curl -X GET http://127.0.0.1:5984
```

```
200 OK
```

```
{  
  "couchdb": "Welcome",  
  "version": "1.0.0"  
}
```

Database

```
curl -X PUT http://127.0.0.1:5984/mydb
```

201 Created

```
{"ok":true}
```

```
curl -X GET http://127.0.0.1:5984/mydb
```

```
200 OK
```

```
{  
  "db_name": "mydb",  
  "doc_count": 0,  
  "update_seq": 0,  
  "compact_running": false,  
  ...  
}
```

```
curl -X DELETE http://127.0.0.1:5984/mydb
```

```
200 OK
```

```
{"ok":true}
```

```
curl -X GET http://127.0.0.1:5984/mydb
```

404 Object Not Found

```
{  
  "error": "not_found",  
  "reason": "no_db_file"  
}
```

Documents

```
curl -X PUT http://127.0.0.1:5984/mydb/mydoc \  
-d '{"yo":"dawg"}'
```

201 Created

```
{  
  "ok":true,  
  "id":"mydoc",  
  "rev":"1-a737e531d8f096d1aa78d9e11704ae50"  
}
```



```
curl -X GET http://127.0.0.1:5984/mydb/mydoc
```

```
200 OK
```

```
{  
  "_id": "mydoc",  
  "_rev": "1-a737e531d8f096d1aa78d9e11704ae50",  
  "yo": "dawg"  
}
```

IDs and Revisions

IDs

- Every document has an ID that is unique for a database
- The ID must be a string
- You can use special characters, but you have to URL escape them
- The default document ID is a UUID

Revisions

- MVCC (Multi Version Concurrency Control)
- Documents aren't updated in place
- Instead a new revision is created
- Does not give you versioning of documents for free!
- Old revisions are removed on compaction

```
curl -X PUT http://127.0.0.1:5984/mydb/mydoc \  
-d '{"hello":"world"}'
```

409 Conflict

```
{  
  "error": "conflict",  
  "reason": "Document update conflict."  
}
```

```
curl -X PUT http://127.0.0.1:5984/mydb/mydoc \  
-d '{"_rev":"1-a737e531d8f096d1aa78d9e11704ae50",  
      "hello":"world"}'
```

201 Created

```
{  
  "ok":true,  
  "id":"mydoc",  
  "rev":"2-2e6854228b269fcb5d1899e894773211"  
}
```

```
curl -X GET http://127.0.0.1:5984/mydb/mydoc
```

```
200 OK
```

```
{  
  "_id": "mydoc",  
  "_rev": "2-2e6854228b269fcb5d1899e894773211",  
  "hello": "world"  
}
```

```
curl -X DELETE http://127.0.0.1:5984/mydb/mydoc?  
rev=2-2e6854228b269fcb5d1899e894773211
```

200 OK

```
{  
  "ok":true,  
  "id":"mydoc",  
  "rev":"3-bdf9cf79f24b1cef35633d57a167023f"  
}
```



```
curl -X GET http://127.0.0.1:5984/mydb/mydoc
```

404 Object Not Found

```
{  
  "error": "not_found",  
  "reason": "deleted"  
}
```

Attachments

```
curl -H 'Content-Type: image/png' \  
      -T image.png -X PUT \  
http://127.0.0.1:5984/mydb/mydoc/gnu.png?rev=1-15a4
```

201 Created

```
{  
  "ok":true,  
  "id":"mydoc",  
  "rev":"2-1822fe627f299844df7412e254364b7b"  
}
```

```
curl -X GET http://127.0.0.1:5984/mydb/mydoc
```

200 OK

```
{  
  "_id": "mydoc", "_rev": "2-1822", "hello": "world",  
  "_attachments": {  
    "gnu.png": {  
      "content_type": "image/png",  
      "revpos": 2,  
      "length": 63573,  
      "stub": true  
    }  
  }  
}
```

```
curl -X GET -O \  
http://127.0.0.1:5984/mydb/mydoc/gnu.png
```

200 OK



Views

Map/Reduce

Map

Map Function

```
function(doc) {  
  emit(doc._id, doc);  
}
```

Map Function

- Incremental
- Runs in isolation
- `emit(key, value)`
- Can call emit 0 to N times
- Results are ordered by key
- Key can be simple (number, string) or complex (array or hash)
- Value can be null, a number, string, array or hash (it is just another document)

Reduce

Reduce Function

```
function(key, values, rereduce) {  
    var sum = 0;  
    for (var i = 0; i < values.length; i++) {  
        sum += values[i];  
    }  
    return sum;  
}
```

Reduce Function

- Optional
- Aggregates the values emitted by Map
- Must be commutative
- $(a + b) = (b + a)$
- Must be associative
- $a + (b + c) = (a + b) + c$
- rereduce - Will be true if values being passed in are the results of a previous reduction

Design Documents

Design Documents

- System document that defines application logic
- The ID begins with _design/

```
curl -X GET http://127.0.0.1:5984/mydb/_design/myapp
```

```
200 OK
```

```
{
  "_id": "_design/myapp",
  "_rev": "1-2732c3d854ac67ec8717227134705623",
  "language": "javascript",
  "views": {
    "myview": {
      "map": "function(doc) { ... }",
      "reduce": "function(key, val, rered) { ... }"
    }
  }
}
```



```
{
  "_id": "Receipt-1",
  "_rev": "1-827",
  "created_on": "2010-09-07",
  "total": 1
}
```

```
{
  "_id": "Receipt-2",
  "_rev": "1-b2a",
  "created_on": "2010-09-01",
  "total": 2
}
```

```
{
  "_id": "Receipt-3",
  "_rev": "1-c92",
  "created_on": "2010-09-01",
  "total": 3
}
```

```
{
  "_id": "Receipt-4",
  "_rev": "1-9ff",
  "created_on": "2010-08-31",
  "total": 4
}
```

```
"views": {  
  "myview": {  
    "map": "function(doc) {  
      emit(doc.created_on.split("-"), doc.total);  
    }"  
  }  
}
```

```
curl -X GET http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview
```

200 OK

```
{"total_rows":4, "offset":0, "rows":  
  {"id":"Receipt-4", "key":["2010", "08", "15"], "value":4},  
  {"id":"Receipt-3", "key":["2010", "09", "01"], "value":3},  
  {"id":"Receipt-2", "key":["2010", "09", "01"], "value":2},  
  {"id":"Receipt-1", "key":["2010", "09", "07"], "value":1}  
]}
```

```
curl -g -X GET 'http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview?key=["2010","09","01"]'
```

200 OK

```
{"total_rows":4, "offset":1, "rows":  
  {"id":"Receipt-3", "key":["2010","09","01"], "value":3},  
  {"id":"Receipt-2", "key":["2010","09","01"], "value":2}  
]}
```

```
curl -g -X GET 'http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview?startkey=["2010","09"]  
&endkey=["2010","09",{ }]'
```

200 OK

```
{"total_rows":4, "offset":1, "rows":  
  {"id":"Receipt-3", "key":["2010","09","01"],"value":3},  
  {"id":"Receipt-2", "key":["2010","09","01"],"value":2},  
  {"id":"Receipt-1", "key":["2010","09","07"],"value":1}  
]}
```

```
"views": {  
  "myview": {  
    "map": "function(doc) {  
      emit(doc.created_on.split("-"), doc.total);  
    }",  
    "reduce": "_sum"  
  }  
}
```

```
curl -X GET http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview
```

200 OK

```
{"rows": [  
  {"key": null, "value": 10}  
]}
```

```
curl -X GET http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview?group=true
```

200 OK

```
{"rows": [  
  {"key": ["2010", "08", "15"], "value": 4},  
  {"key": ["2010", "09", "01"], "value": 5},  
  {"key": ["2010", "09", "07"], "value": 1}  
]}
```



```
curl -X GET http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview?group_level=2
```

200 OK

```
{"rows": [  
  {"key": ["2010", "08"], "value": 4},  
  {"key": ["2010", "09"], "value": 6}  
]}
```

```
curl -X GET http://127.0.0.1:5984/mydb/_design/  
myapp/_view/myview?group_level=1
```

200 OK

```
{"rows": [  
  {"key": ["2010"], "value": 10}  
]}
```

View Query Options

- key
- startkey
- startkey_docid
- endkey
- endkey_docid
- limit
- stale
- descending
- skip
- group
- group_level
- reduce
- include_docs
- inclusive_end

Other Design Document Functions

- shows
- lists
- validate_doc_update
- filters

Indexers

Lucene
(full text search)

Geocouch
(geospatial queries)

Replication

```
curl -X POST http://127.0.0.1:5984/_replicate \
-d '{"source":"mydb","target":"http://me.com/mydb"}'
```

200 OK

```
{"ok":true,
 "session_id":"bcc26cc1cbe708fc8f2d72c7cf831181",
 "source_last_seq":30,
 "history":[
  {"session_id":"bcc26cc1cbe708fc8f2d72c7cf831181",
   "start_time":"Tue, 28 Sep 2010 18:09:57 GMT",
   "end_time":"Tue, 28 Sep 2010 18:09:59 GMT",
   "start_last_seq":0,"end_last_seq":30,"recorded_seq":30,
   "missing_checked":0,"missing_found":6,
   "docs_read":7,"docs_written":6,"doc_write_failures":1
  }
]
```


Replication

- source and target can be local or remote
- {"continuous":true}
- If a conflict occurs during replication, a winner is deterministically chosen
- The loser is saved with the doc and can be accessed with ?conflicts=true

Replication Filters

```
{"filter": "myapp/myfilter"}

"filters": {
  "myfilter": {
    "function(doc, req) {
      if (doc.type == 'Post' &&
          req.userCtx.name == doc.author) {
        return true;
      } else {
        return false;
      }
    }
  }
}
```

`_changes`

```
curl -X GET http://127.0.0.1:5984/mydb/_changes
```

200 OK

```
{"results": [
  {"seq": 1, "id": "doc1",
    "changes": [{"rev": "1-42dd"}]},
  {"seq": 2, "id": "doc1",
    "changes": [{"rev": "2-9bf5"}]},
  {"seq": 3, "id": "doc2",
    "changes": [{"rev": "1-737d"}]},
  {"seq": 4, "id": "doc1",
    "changes": [{"rev": "3-8822"}], "deleted": true},
  "last_seq": 4}
```

```
curl -X GET http://127.0.0.1:5984/mydb/_changes?  
since=2
```

200 OK

```
{"results": [  
  {"seq": 3, "id": "doc2",  
    "changes": [{"rev": "1-737d"}]},  
  {"seq": 4, "id": "doc1",  
    "changes": [{"rev": "3-8822"}], "deleted": true},  
  "last_seq": 4}
```

```
curl -X GET http://127.0.0.1:5984/mydb/_changes?  
feed=continuous
```

200 OK

```
{"seq":1,"id":"doc1","changes":[{"rev":"1-42dd"}]}}  
{"seq":2,"id":"doc1","changes":[{"rev":"2-9bf5"}]}}  
{"seq":3,"id":"doc2","changes":[{"rev":"1-737d"}]}}  
{"seq":4,"id":"doc1",  
  "changes":[{"rev":"3-8822"}],"deleted":true}}
```

...

```
{"seq":5,"id":"doc3","changes":[{"rev":"1-b12c"}]}}
```

_changes Options

- since
- feed (normal | longpoll | continuous)
- heartbeat
- timeout
- filter
- include_docs

Security

Admin Party

by default, everyone is an admin

Admin Rights

- Create and delete databases
- Create, update and delete design documents
- Trigger compaction
- Reading the task status list
- Restart the server
- Read and update the active configuration

Authentication

- HTTP Basic
- Secure Cookie
- OAuth

Authorization

- Can define admins for a particular database
- The admins can update design documents and create database readers
- Can define readers for a particular database
- Readers can read documents
- Database is public if no readers are defined
- Can define list of roles for admins and readers

Authorization

- Users stored in `_users` database
- Can use `validate_doc_update` to reject an update based on the currently logged in user

Sharding/Clustering

Lounge

(NGINX and python twisted daemon)

BigCouch

(Erlang)

Ruby Clients

Net::HTTP

RestClient

couchrest
couch_potato
couch_foo
relaxdb

couchrest_extended_document
Rails 2.x

couchrest_model
Rails 3.0

```
class Page < CouchRest::Model::Base
  property :slug, String
  property :body, String
  property :authors, [String]
  timestamps!

  validates_presence_of :slug

  view_by :slug
end
```

Negative

- Slow compared to other NoSQL solutions
- CouchDB is a glutton for disk space
- Run compaction often
- Can't chain Map/Reduce (yet)

CouchApp

Bind your Couch to
your public IP address